

# MNIST Digit Classification with k-NN and Decision Trees

Ada Boran Yilmaz

# 1 Overview

This report presents a comprehensive approach to classifying handwritten digits (0–9) from the MNIST dataset using two different machine learning algorithms:

- k-Nearest Neighbors (k-NN)
- Decision Tree

We begin with an overview of the dataset and our preprocessing steps. Next, we detail the hyperparameter tuning for each model and present the final evaluation results on a held-out test set. We conclude with an analysis of misclassifications and potential explanations for errors.

# 2 Dataset and Preprocessing

## 2.1 Data Loading

- Loads the MNIST data using Keras API
- Splits the training data as 80% training set and 20% validation set
- Verifies the shapes

Training set shape: (48000, 28, 28)

Validation set shape: (12000, 28, 28)

Test set shape: (10000, 28, 28)

# 2 Dataset and Preprocessing


## 2.2 Data Analysis

- Computes class distribution for training data
- Plots class distribution as a bar chart
- Calculates the basic statistics such as mean and standard deviation of pixel values.
- Visualizes one sample per digit

Class Distribution: {0: 4729, 1: 5470, 2: 4762, 3: 4889, 4: 4655, 5: 4324, 6: 4748, 7: 4968, 8: 4653, 9: 4802}

# 2 Dataset and Preprocessing

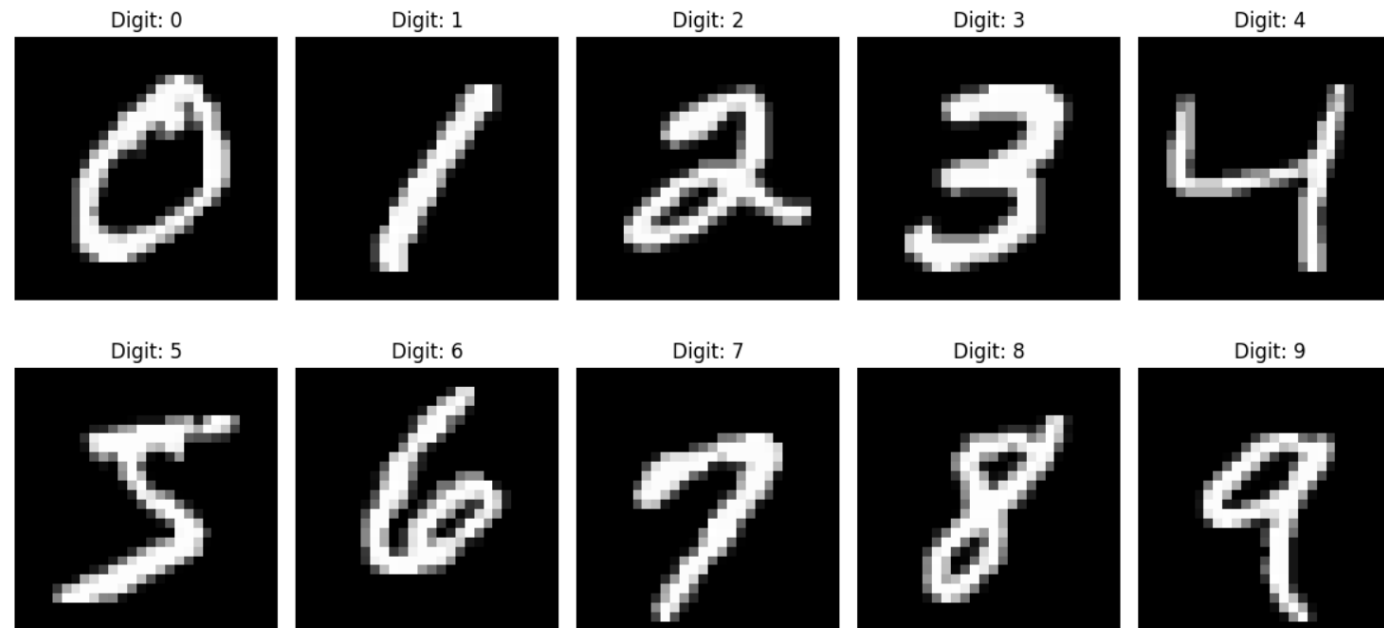
## 2.2 Data Analysis



# 2 Dataset and Preprocessing

## 2.2 Data Analysis

Mean pixel value: 33.37, Standard Deviation of pixel values: 78.64



# 2 Dataset and Preprocessing

## 2.3 Data Preprocessing

- Normalizes the images by scaling pixel values to  $[0, 1]$
- Verifies that the shapes are unchanged

Training set shape (normalized): (48000, 28, 28)

Validation set shape (normalized): (12000, 28, 28)

Test set shape (normalized): (10000, 28, 28)

# 3 k-NN Classifier

## 3.1 Model Initialization and Hyperparameter Tuning

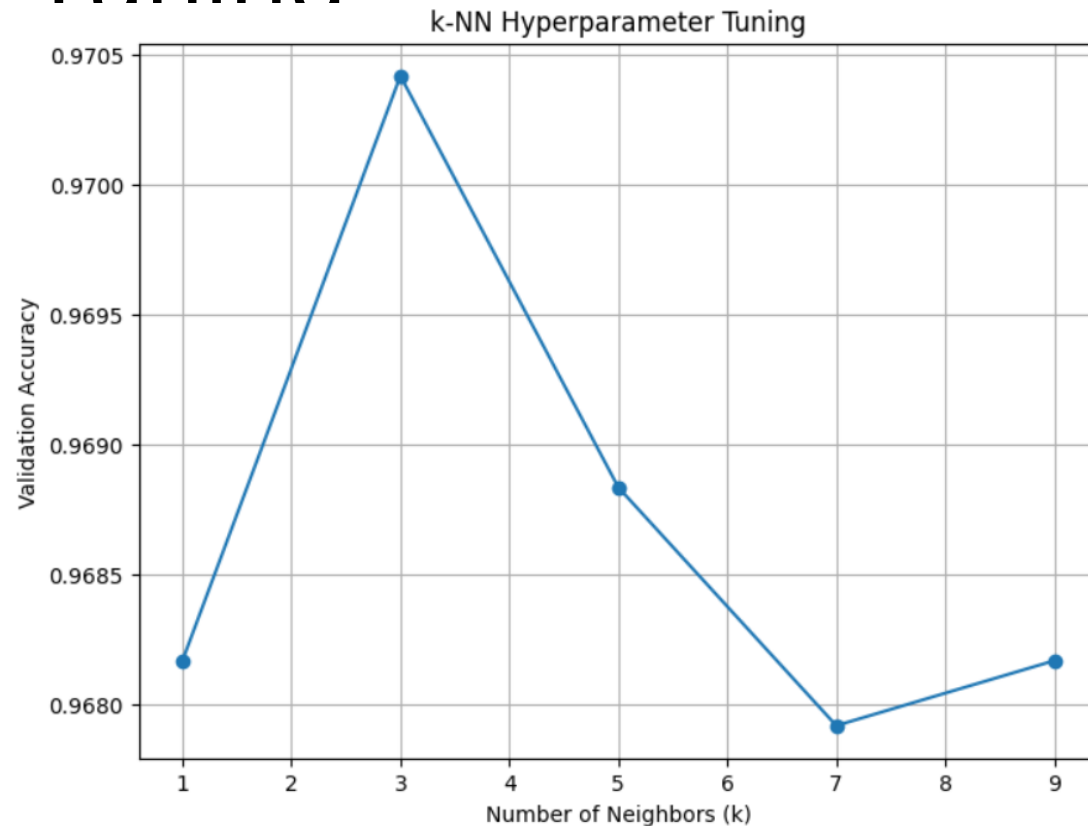
- Flattens the 28x28 images for training and validation
- Defines the neighbor values to try
- Loops over each neighbor value and evaluate on the validation set
- Plots the validation accuracy versus the number of neighbors

k = 1, Validation Accuracy: 0.9682; k = 3, Validation Accuracy: 0.9704;  
k = 5, Validation Accuracy: 0.9688; k = 7, Validation Accuracy: 0.9679;  
k = 9, Validation Accuracy: 0.9682



# 3 k-NN Classifier

## 3.1 Model Initialization and Hyperparameter Tuning



# 3 k-NN Classifier

## 3.2 Final Model Training and Evaluation

- Concatenates training and validation sets, then reshapes images into 2D arrays.
- Uses the best k from validation to fit the final k-NN classifier.
- Predicts labels on the test set and computes accuracy, precision, recall, and F1-score.
- Displays a confusion matrix and calculates per-digit misclassifications.
- Identifies and prints the top three digits with the highest misclassification rates.
- Randomly selects and visualizes five misclassified examples with true and predicted labels.

# 3 k-NN Classifier

## 3.2 Final Model Training and Evaluation

Optimal number of neighbors (k): 3

Test Accuracy: 0.9705


Test Precision: 0.9707

Test Recall: 0.9705

Test F1 Score: 0.9705

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	980
1	0.96	1.00	0.98	1135
2	0.98	0.97	0.97	1032
3	0.96	0.97	0.96	1010
4	0.98	0.97	0.97	982
5	0.97	0.96	0.96	892
6	0.98	0.99	0.98	958
7	0.96	0.96	0.96	1028
8	0.99	0.94	0.96	974
9	0.96	0.96	0.96	1009
accuracy			0.97	10000
macro avg	0.97	0.97	0.97	10000
weighted avg	0.97	0.97	0.97	10000



# 3 k-NN Classifier

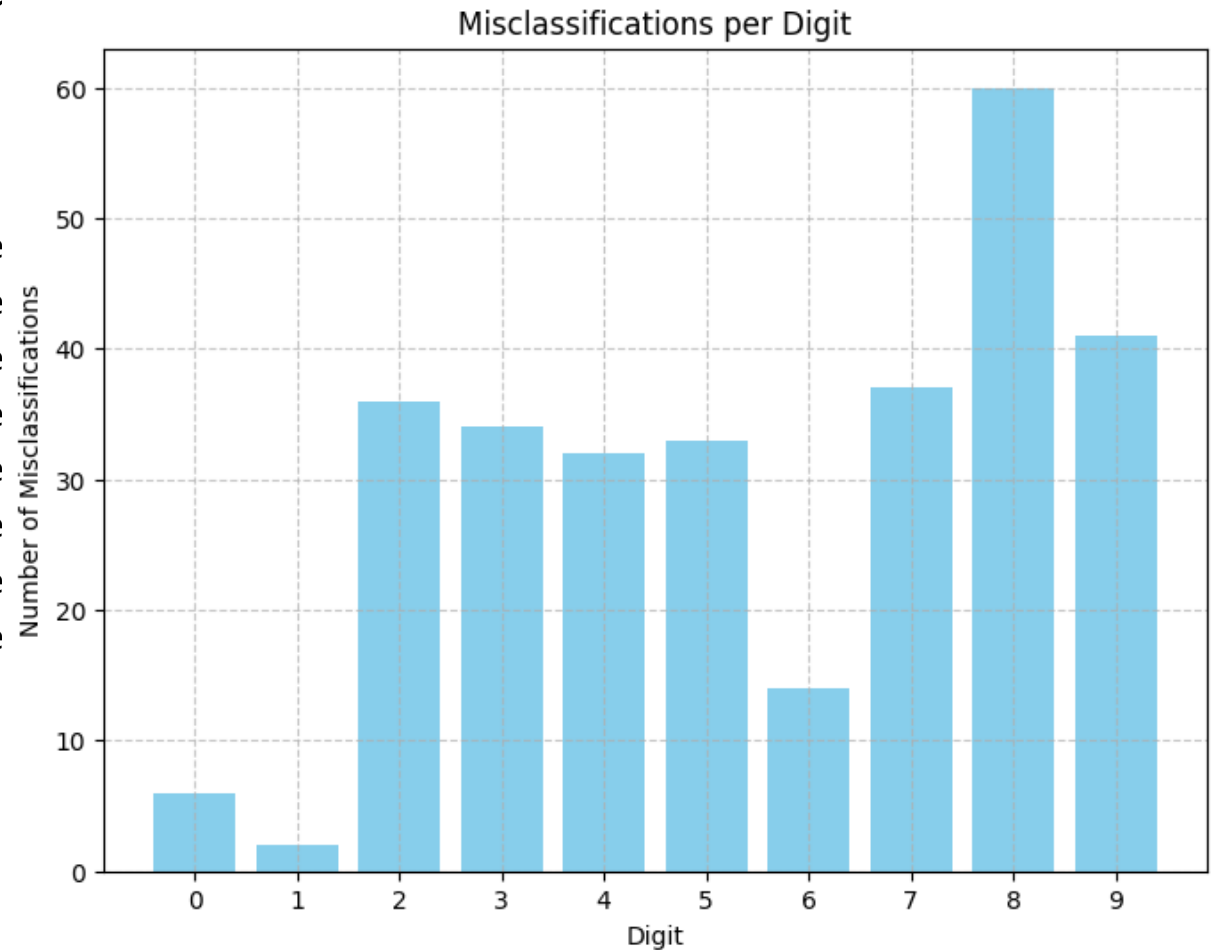
## 3.2 Final Model Training and Evaluation

Total misclassified examples: 295

- Digit 0: 6 misclassifications, 0.61% misclassification rate
- Digit 1: 2 misclassifications, 0.18% misclassification rate
- Digit 2: 36 misclassifications, 3.49% misclassification rate
- Digit 3: 34 misclassifications, 3.37% misclassification rate
- Digit 4: 32 misclassifications, 3.26% misclassification rate
- Digit 5: 33 misclassifications, 3.70% misclassification rate
- Digit 6: 14 misclassifications, 1.46% misclassification rate
- Digit 7: 37 misclassifications, 3.60% misclassification rate
- Digit 8: 60 misclassifications, 6.16% misclassification rate
- Digit 9: 41 misclassifications, 4.06% misclassification rate

Top 3 Most Frequently Misclassified Digits:

- Digit 8: 60 misclassifications, Rate: 6.16%
- Digit 9: 41 misclassifications, Rate: 4.06%
- Digit 7: 37 misclassifications, Rate: 3.60%




# 3 k-NN Classifier

## 3.2 Final Model Training and Evaluation

5 Random Misclassified Examples:


True: 3  
Pred: 5




True: 2  
Pred: 0




True: 8  
Pred: 0



True: 7  
Pred: 1



True: 6  
Pred: 4



# 4 Decision Tree Classifier

## 4.1 Model Training and Hyperparameter Tuning

- Reshapes the training and validation sets into 2D arrays so that the Decision Tree classifier can process the data. Defines the grid of hyperparameter values for `max_depth` and `min_samples_split` to explore during tuning.
- Creates a dictionary to store the validation accuracy for each hyperparameter combination and initializes variables to track the best configuration.
- Iterates over every combination of `max_depth` and `min_samples_split`, training a Decision Tree for each setting. For each configuration, initializes the classifier with the current hyperparameters, fits it on the flattened training data, and predicts on the validation set.
- Computes the validation accuracy for the current configuration and stores the result, printing the configuration and its accuracy.
- Updates the best hyperparameters if the current configuration yields a higher validation accuracy. Prints out the best-performing hyperparameter combination along with its corresponding accuracy.

# 4 Decision Tree Classifier

## 4.1 Model Training and Hyperparameter Tuning

Max Depth: 2, Min Samples Split: 2 => Validation Accuracy: 0.1872

Max Depth: 2, Min Samples Split: 5 => Validation Accuracy: 0.1872

Max Depth: 5, Min Samples Split: 2 => Validation Accuracy: 0.2855

Max Depth: 5, Min Samples Split: 5 => Validation Accuracy: 0.2855

Max Depth: 10, Min Samples Split: 2 => Validation Accuracy: 0.4176

Max Depth: 10, Min Samples Split: 5 => Validation Accuracy: 0.4177

Best Hyperparameters:

Max Depth: 10, Min Samples Split: 5, Accuracy: 0.4177

# 4 Decision Tree Classifier

## 4.2 Evaluation

- Combines the training and validation sets and reshapes the images into 2D arrays so the classifier can process them.
- Initializes the final Decision Tree classifier with the best hyperparameters (max\_depth=10, min\_samples\_split=5) and trains it on the combined dataset.
- Predicts the labels on the test set and computes evaluation metrics including accuracy, precision, recall, and F1-score, then prints a detailed classification report.
- Generates and visualizes a confusion matrix using a heatmap to analyze which classes are most often misclassified.
- Provides a discussion prompt to examine the confusion matrix for patterns in misclassification (e.g., identifying digits that are frequently confused).
- Computes predicted probabilities for each class, binarizes the test labels for a multi-class ROC analysis, calculates ROC curves and AUC scores for each digit, and plots all ROC curves on a single graph with AUC values included in the legend.




# 4 Decision Tree Classifier

## 4.2 Evaluation

Test Accuracy: 0.3714  
Test Precision: 0.6540  
Test Recall: 0.3714  
Test F1 Score: 0.3609

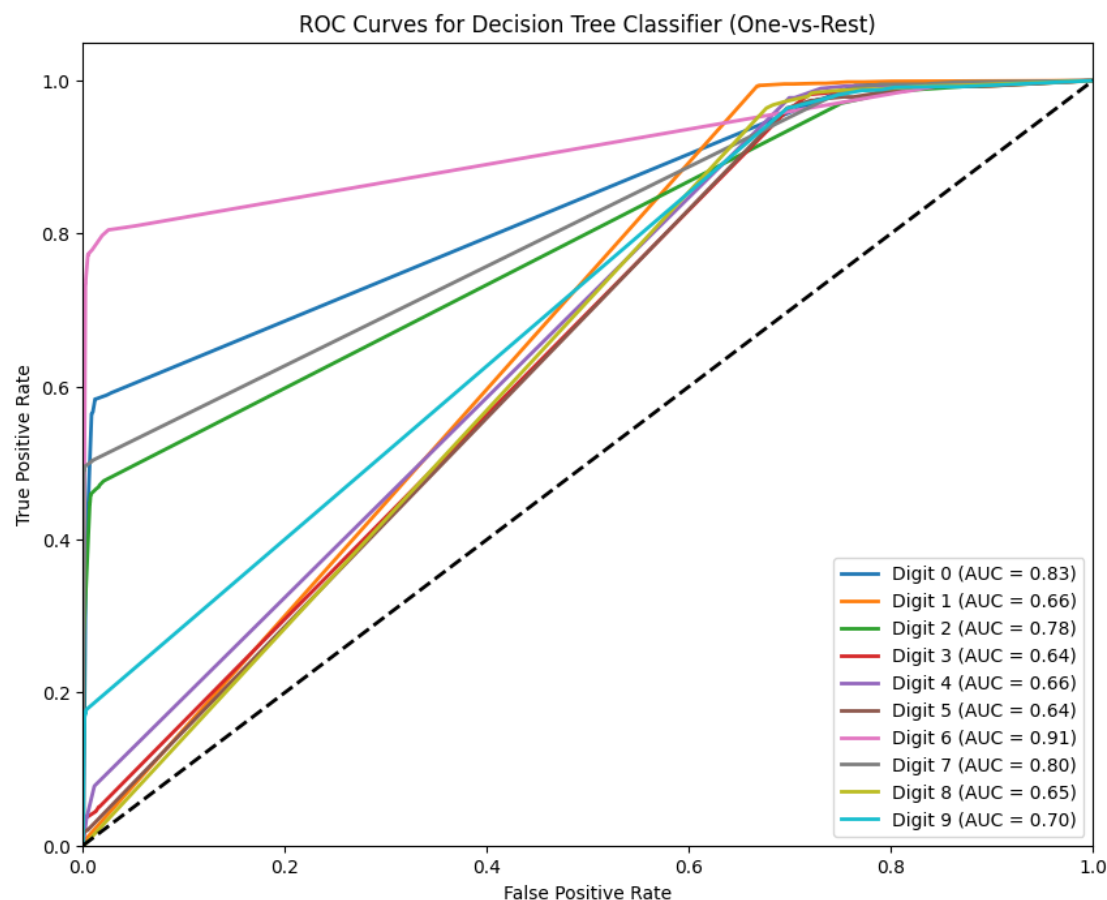
Detailed Classification Report:

	precision	recall	f1-score	support
0	0.86	0.58	0.69	980
1	0.16	0.99	0.28	1135
2	0.87	0.46	0.60	1032
3	0.73	0.03	0.07	1010
4	0.43	0.08	0.13	982
5	0.74	0.02	0.03	892
6	0.94	0.77	0.85	958
7	0.97	0.50	0.65	1028
8	0.00	0.00	0.00	974
9	0.90	0.17	0.29	1009
accuracy			0.37	10000
macro avg	0.66	0.36	0.36	10000
weighted avg	0.65	0.37	0.36	10000



# 4 Decision Tree Classifier

## 4.2 Evaluation



# 4 Decision Tree Classifier

## 4.2 Evaluation

Decision Tree Confusions:

- Almost all numbers other than 6 are confused for number 1.
- 0, 1, 6 and 7 are the only numbers that has the top predictions in their own category.
- Number 3, 4, 5, 8 and 9 all were predicted as mostly 1.